

Transformers II

Transformers

- We will continue to look at
 - Encoder-only/decoder-only models
 - In-context Learning
 - Parameter-efficient transformers
 - Low-rank adapters
 - Model distillation
 - Sentence Transformers
 - Reinforcement learning with transformers
-

Encoder-Decoder Transformers

- Originally, transformers had an encoder and decoder
 - The encoder for input sequence and decoder for output sequence
 - For machine translation, the decoder produces the text in the foreign language
 - For language modelling, the decoder produces the next word in the sequence
 - Decoder output is input shifted by one word
-

Encoder-Only Transformers

- In encoder-only transformers, the decoder is removed
 - These models produce an output for the sequence
 - For example, in BERT the output is a vector for each token in the input sequence
-

Decoder-Only Transformers

- In decoder-only transformers, the encoder is removed
 - This is used by models such as GPT
 - They are focused on generation and trained on next word prediction
-

Fine-tuning Transformers

- Transformers are typically trained on a large corpus of text

- The model is then fine-tuned on a specific task
 - In order to fine-tune, we take the model and then add a task-specific head
 - The model is trained as usual, but the weights are initialized from the pre-trained model
-

In-content Learning

- In-context learning is a technique to adapt a pre-trained model to a specific task
 - It does not require fine-tuning the entire model (i.e., updating the weights)
 - Transformers are capable of in-context learning if trained on a diverse set of tasks
-

Prompt engineering

- Techniques for prompt engineering:
 - Zero-shot learning
 - Few-shot learning
 - Chain of thought
 - Prefix tuning
-

Parameter-efficient Transformers

- Transformers are very large models
 - Most don't fit on a single GPU
 - Reducing the number of parameters is possible
 - Trade-off between model size and performance
-

Model distillation

- Model distillation is a technique to train a smaller model to mimic a larger model
 - The smaller model is trained on the output of the larger model
 - DistilBERT was able to reduce the size of BERT by 40% while maintaining 97% of the performance
-

Parameter-efficient Fine-tuning

- Language models such as GPT-3 have billions of parameters

- This makes them very expensive to train
 - Parameter-efficient transformers are designed to be smaller and faster to train
-

Low-Rank Adapters (LoRA)

- Low-Rank Adapters (LoRA) are a technique to allow efficient training of large models
- The idea is to add a small adapter layer to each transformer layer
- This adapter layer calculates a low-rank delta, ΔW , that is added to the weights of the self-attention layer

$$W_{new} = W_{old} + \Delta W$$




Figure 1: Low-Rank Adapters

Low-Rank Adapters (LoRA)

- The delta is a product of two low-rank matrices with $r \ll d$.
 - Faster to train
 - Smaller memory footprint
 - Not less than inference on the full model!
-

Quantization

- Quantization is a technique to reduce the precision of the weights
 - Instead of 32-bit floating-point numbers, we use 8-bit integers
 - This reduces the memory footprint and speeds up training
-

Quantization Precision Levels

- Floating-point numbers consist of a sign bit, exponent, and mantissa

$$\text{float} = (-1)^s \times 2^e \times 1.m$$

NB Mantissa is a binary fraction

Quantization Precision Levels

Bits	Sign	Exponent	Mantissa
64	1	11	52
32	1	8	23
16	1	5	10
8	1	3	4

Affine Quantization

We can quantize to integers using the following formula:

$$x_q = \text{round}\left(\frac{x}{s} + Z\right)$$

Where s is the scale factor and Z is the zero-point

This allows us to quantize down to 8, 4 or even 2 bits.

Quantization

- Quantization can be done in two ways:
 - Post-training quantization
 - Quantization-aware training
-

Post-training Quantization

- Post-training quantization is done after the model is trained
 - The weights are quantized to a lower precision
 - Further training on the quantized model can improve performance
-

Quantization-aware Training

- During forward passes, weights and activations are simulated as quantized (e.g., 8-bit integers)
 - Stored as floating-point numbers during training.
 - Gradients are calculated in full precision.
 - Enables higher accuracy and fast inference.
-

Sentence Transformers

- Sentence transformers are models that take a sentence as input and produce a vector representation
 - These vectors can be used for tasks such as semantic similarity, clustering, and classification
-

Siamese Networks

- Siamese network, uses two similar examples x_1 and x_2 .

$$h_1 = f(x_1)$$

$$h_2 = f(x_2)$$

$$\text{similarity} = \text{cosine}(h_1, h_2)$$

Triplet loss

- Triplet loss, uses a positive example x_p and a negative example x_n .

$$\text{loss} = \max(0, \text{margin} - \text{cosine}(h_a, h_p) + \text{cosine}(h_a, h_n))$$

Sentence Transformer Output

- Mean pooling: average of the token embeddings
 - Max pooling: maximum of the token embeddings
 - CLS token: the output of a special token [CLS] in the transformer
-

Using Sentence Transformers

- Textual similarity: cosine similarity between two vectors
 - Classification: add a new network to predict class from vector
 - Clustering: use k-means or hierarchical clustering
 - Information Retrieval: use vectors to retrieve similar documents; often using vector database such as FAISS
-

Reinforcement Learning with Transformers

- Combines reinforcement learning with human-provided feedback
 - Used to train models such as GPT
-

RLHF

1. Pre-train a model on a large corpus
2. Collect human feedback
 - Model generates answers
 - Human annotators rank results

3. Train a reward model
 - Predict a numerical reward from human feedback
 4. Fine-tune the model using reinforcement learning
 - Use Proximal Policy Optimization (PPO) or other RL algorithms
-

Why Reinforcement Learning?

- Align model with human values
 - Handle ambiguous or subjective tasks
 - Mitigate harmful behaviour
-

Summary

- Transformers are a powerful architecture for NLP
- Most transformers are encoder-only or decoder-only
- Transformers can be fine-tuned for specific tasks
- Parameter efficient transformers can be trained more easily
- Sentence transformers are used for semantic similarity and other tasks
- Reinforcement learning can make transformers more human-like